

```
In [3]: import os

gen_dog_imgs = '/kaggle/working/generative-dog-images'
if not os.path.exists(img_dog_imgs):
    os.makedirs(gen_dog_imgs)

dogs_dir = '/kaggle/working/dogs'
if not os.path.exists(dogs_dir):
    os.makedirs(dogs_dir)

# mkdir /kaggle/working/generative-dog-images
!mkdir /kaggle/working/generative-dog-images/all-dogs.zip -d /kaggle/working/generative-dog-images > /dev/null 2>&1
!unzip /kaggle/input/generative-dog-images/Annotation.zip -d /kaggle/working/generative-dog-images > /dev/null 2>&1
```

## Generative Adversarial Network

Generative Adversarial Network (GAN) is a class of machine learning frameworks and a prominent framework for approaching generative AI. In a GAN, two neural networks, generator and discriminator, contest with each other in the form of a zero-sum game, where one agent's gain is another agent's loss. The core idea of a GAN is to train the generator to "fool" the discriminator rather than directly minimize the individual image distances, and the discriminator is indirectly trained to tell how realistic the generated images may seem. This way, both the generator and the discriminator are updated dynamically against each other to achieve realistic imitation to the original images.

GAN aims to learn to generate new data with the same statistics as the provided training set. A GAN trained on photographs can generate new images superficially authentic to human observers. While GAN is originally intended to be implemented in unsupervised learning, variations of GANs are developed into models suitable for semi-supervised and supervised learning purposes.

## Dataset Overview

The Stanford Dogs dataset contains images of 120 breeds of dogs worldwide. This dataset has been built using images and annotation from ImageNet for the task of fine-grained image categorization. There are 20,580 images, out of which 12,000 are used for training and 8580 for testing. Class labels and bounding box annotations are provided for all the 12,000 images. In this study, training and testing images are not distinguished as they are repurposed into image generation.

```
In [3]: #matplotlib inline
import torch
import torch.nn as nn
from torchvision import transforms
from torchvision.utils import make_grid
from torch.utils.data import Dataset, DataLoader
import matplotlib.pyplot as plt
import random
from PIL import Image
import xml.etree.ElementTree as ET

In [3]: # Set random seed for reproducibility
manualSeed = 999
print("Random Seed: ", manualSeed)
random.seed(manualSeed)
torch.manual_seed(manualSeed)
torch.set_deterministic(True) # Needed for reproducible results

Random Seed: 999

In [4]: all_dogs_dir = '/kaggle/working/generative-dog-images/all-dogs'
annotation_dir = '/kaggle/working/generative-dog-images/Annotation'

In [5]: def bndbox_extraction(filename, square=False):
    root = ET.parse(filename).getroot()
    box = root.find('object/bndbox')
    xmin, ymin, xmax, ymax = map(int, [box.find(tag).text for tag in ('xmin', 'ymin', 'xmax', 'ymax')])

    if square:
        center_x, center_y = (xmin + xmax) // 2, (ymin + ymax) // 2
        size = max(xmax - xmin, ymax - ymin)
        xmin, xmax = center_x - size // 2, center_x + size // 2
        ymin, ymax = center_y - size // 2, center_y + size // 2

    return xmin, ymin, xmax, ymax
```

The purpose of excluding the "intruders" images in a Dog Image Generation project using DCGAN (Deep Convolutional Generative Adversarial Network) can be explained as follows:

Consistency and Purity of Dataset: By excluding images labeled as "intruders," you ensure that your training data consists only of clear, relevant examples of what you want your model to learn - in this case, dogs. This helps in maintaining the purity of the dataset, focusing solely on the intended subject matter.

Reducing Noise and Improving Model Learning: "Intruders" might include images that are not dogs or are of poor quality, mislabeled, or in some way not representative of the target class (dogs). Including these could introduce noise, making it harder for the model to learn the specific features of dogs. This exclusion helps in reducing such noise, allowing the model to learn more accurate and detailed features of dogs.

Enhancing Model Performance: A cleaner dataset leads to better performance. By removing irrelevant or ambiguous images, you increase the likelihood that the generator and discriminator in the DCGAN will learn to produce and identify dog images with higher fidelity and accuracy. This also potentially reduces the training time as the model doesn't have to waste computational resources on learning from or distinguishing irrelevant images.

Ref: <https://www.kaggle.com/korova/dogs-images-intruders-extraction>

```
In [6]: intruders = [
    'n02088238_10870_0',
    'n02088466_6901_1', 'n02088466_6963_0', 'n02088466_9167_0',
    'n02088466_9167_1', 'n02088466_9167_2',
    'n02088995_2221_0', 'n02088995_2227_1',
    'n02089973_1132_0', 'n02089973_1385_3', 'n02089973_1458_1',
    'n02089973_1799_2', 'n02089973_2793_3', 'n02089973_4055_0',
    'n02089973_4182_1', 'n02089973_4183_0',
    'n02090379_4673_1', 'n02090379_4870_1',
    'n02090622_7705_1', 'n02090622_9356_1', 'n02090622_9883_1',
    'n02090721_209_1', 'n02090721_1222_1', 'n02090721_1558_1',
    'n02090721_1835_1', 'n02090721_3999_1', 'n02090721_4089_1',
    'n02090721_4096_1',
    'n02091032_722_1', 'n02091032_745_1', 'n02091032_1773_0',
    'n02091032_3992_0',
    'n02091134_2148_1', 'n02091134_34246_2',
    'n02091244_583_1', 'n02091244_2407_0', 'n02091244_3498_1',
    'n02091244_3439_1', 'n02091244_3693_2',
    'n02091467_473_0', 'n02091467_4386_1', 'n02091467_4427_1',
    'n02091467_4558_1', 'n02091467_4560_1',
    'n02091635_1192_1', 'n02091635_1492_0',
    'n02091831_1094_1', 'n02091831_2880_0', 'n02091831_7237_1',
    'n02092002_1931_1', 'n02092002_1937_1', 'n02092002_4218_0',
    'n02092002_4596_0', 'n02092002_5246_1', 'n02092002_5918_0',
    'n02093256_1826_1', 'n02093256_4397_0', 'n02093256_4919_0',
    'n02093428_5462_0', 'n02093428_6949_1'
]

Intruders = [filename + '.jpg' if not filename.endswith('.svg') else filename for filename in intruders]
len(intruders)

Out[6]: 60
```

## Compare the number of images in the dataset before and after the intruder exclusion

```
In [7]: dogs_count = 0
for breed in os.listdir(annotation_dir):
    for dog in os.listdir(os.path.join(annotation_dir, breed)):
        #print(dog)
        bndbox = bndbox_extraction(os.path.join(annotation_dir, breed, dog), square=True)
        jpg_name = os.path.join(all_dogs_dir, dog + '.jpg')
        intruders_join = '\t'.join(intruders)
        if os.path.exists(jpg_name):
            if dog not in intruders_join:
                img = Image.open(jpg_name).crop(bndbox)
                img.save(os.path.join(dogs_dir, dog + '.jpg'))
                dogs_count += 1

print('number of dogs in the original dataset:', len(os.listdir(dogs_dir)) - len(intruders))
print('number of dogs in the dataset excluding intruders:', dogs_count)

number of dogs in the original dataset: 20580
number of dogs in the dataset excluding intruders: 20520
```

```
In [8]: class LoadDataset(Dataset):
    def __init__(self, data_dir, transforms=None):
        self.files = os.path.join(data_dir, file) for file in os.listdir(data_dir))
        self.transforms = transforms

    def __len__(self):
        return len(self.files)

    def __getitem__(self, index):
        img = Image.open(self.files[index])
        if self.transforms is not None:
            img = self.transforms(img)
        return img

In [9]: img_size = (64, 64)
img_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Resize(img_size),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

batch_size = 128
trainloader = DataLoader(
    LoadDataset(data_dir=dogs_dir, transform=img_transforms),
    batch_size=batch_size,
    shuffle=True,
    num_workers=2,
)
```

## DCGAN

DCGAN (Deep Convolutional GAN) is a generative adversarial network architecture using deep convolutional neural networks. It is specialized for generating realistic images, mainly square images, in computer vision field. DCGAN can learn and capture detailed features in images of the original training dataset to generate realistic fake images hardly distinguishable by human eyes.

A great DCGAN tutorial can be founded in: [https://pytorch.org/tutorials/beginner/dcgan\\_faces\\_tutorial.html](https://pytorch.org/tutorials/beginner/dcgan_faces_tutorial.html)

```
In [10]: class Generator(nn.Module):
    def __init__(self, z_channels, out_channels=3):
        super(Generator, self).__init__()

        convs = []
        channels = [z_channels, 1024, 512, 256, 128, 64]
        for i in range(1, len(channels)):
            convs.append(nn.ConvTranspose2d(channels[i-1], channels[i], 2, stride=2, bias=False))
            convs.append(nn.BatchNorm2d(channels[i]))
            convs.append(nn.LeakyReLU(0.1, inplace=True))
            convs.append(nn.ConvTranspose2d(channels[i-1], out_channels, 2, stride=2, bias=False))
            convs.append(nn.Tanh())

        self.convs = nn.Sequential(*convs)

    def forward(self, x):
        return self.convs(x)

In [11]: class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()
        channels = [3, 64, 128, 256, 512]
        convs = []
        for i in range(1, len(channels)):
            convs.append(nn.Conv2d(channels[i-1], channels[i], 3, padding=1, stride=2, bias=False))
            if i != 1:
                convs.append(nn.BatchNorm2d(channels[i]))
            convs.append(nn.LeakyReLU(0.2, inplace=True))

        convs.append(nn.Conv2d(channels[-1], 1, 4, bias=False))
        convs.append(nn.Sigmoid())

        self.convs = nn.Sequential(*convs)

    def forward(self, x):
        x = self.convs(x)
        return x.view(-1)

In [12]: def weights_init(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1:
        nn.init.normal_(m.weight.data, 0.0, 0.02)
    elif classname.find('BatchNorm') != -1:
        nn.init.constant_(m.weight.data, 1.0, 0.02)
        nn.init.constant_(m.bias.data, 0)

In [13]: z_channels = 100
G = Generator(z_channels, 3)
G.apply(weights_init)
D = Discriminator()
D.apply(weights_init)
criterion = nn.BCELoss()

cuda = torch.cuda.is_available()
if cuda:
    print("Use GPU")
    G = G.cuda()
    D = D.cuda()
    criterion = criterion.cuda()
else:
    print("No GPU")

Use GPU

In [14]: lr = 0.0004
b1,b2 = 0.4,0.999

optimizerG = torch.optim.Adam(G.parameters())
optimizerD = torch.optim.Adam(D.parameters())

In [15]: fixed_noise = torch.normal(0, 0.1, size=(64, z_channels, 1, 1))
if cuda:
    fixed_noise = fixed_noise.cuda()

epochs = 20
generate_imgs = []
G_losses, D_losses = [], []
for epoch in range(epochs):
    for i, img in enumerate(trainloader):
        z = torch.normal(0, 0.1, size=(img.size(0), z_channels, 1, 1))
        real = torch.ones(img.size(0))
        fake = torch.zeros(img.size(0))
        if cuda:
            img, z = img.cuda(), z.cuda()
            real, fake = real.cuda(), fake.cuda()

        # train D
        D.zero_grad()
        loss_real = criterion(D(img), real)
        loss_real.backward()

        fake_img = G(z)
        loss_fake = criterion(D(fake_img).detach(), fake)
        loss_fake.backward()
        loss_D = (loss_real + loss_fake) / 2


        optimizerD.step()

        # train G
        G.zero_grad()
        loss_G = criterion(D(fake_img), real)
        loss_G.backward()
        optimizerG.step()

        with torch.no_grad():
            noise_img = G(fixed_noise)
            generate_imgs.append(noise_img)
            print(f"[Epoch {epoch+1}]/[epoches] | D loss: {loss_D.item()} | D loss: {loss_D.item()} | loss_real: {loss_real.item()} | loss_fake: {loss_fake.item()}")
            G_losses.append(loss_D.item())
            D_losses.append(loss_D.item())

[Epoch 1/20] | D loss: 10.865918184997559 | D loss: 0.2619585967184448 | loss_real: 0.05483179458379745 | loss_fake: 0.46909973025321861
[Epoch 2/20] | D loss: 6.44848459245981 | D loss: 0.09412280470120202 | loss_real: 0.020279213093136795 | loss_fake: 0.132663993592810061
[Epoch 3/20] | D loss: 4.126104354858398 | D loss: 0.2906491458415985 | loss_real: 0.320544602556854 | loss_fake: 0.25626368952976116
[Epoch 4/20] | D loss: 5.791051131286621 | D loss: 0.1359787583531153 | loss_real: 0.051664398926284866 | loss_fake: 0.2162095737989713
[Epoch 5/20] | D loss: 4.949713171814899 | D loss: 0.0460836042123204 | loss_real: 0.057046431252749 | loss_fake: 0.05127281991958642
[Epoch 6/20] | D loss: 6.157948481579879 | D loss: 0.071864552795887 | loss_real: 0.131789430975914 | loss_fake: 0.0119367740892771
[Epoch 7/20] | D loss: 7.785643861335453 | D loss: 0.18688279390335083 | loss_real: 0.07670797621516436 | loss_fake: 0.29705846040024451
[Epoch 8/20] | D loss: 4.8096537590026855 | D loss: 0.1566897282137148 | loss_real: 0.1338317452449799 | loss_fake: 0.17554774808409294
[Epoch 9/20] | D loss: 6.220575321939678 | D loss: 0.0963636563320159 | loss_real: 0.12573101823149872 | loss_fake: 0.07759813142532361
[Epoch 10/20] | D loss: 7.131747185861511 | D loss: 0.04842083931383054 | loss_real: 0.0416156885823395 | loss_fake: 0.05543811146752739
[Epoch 11/20] | D loss: 9.025453567504883 | D loss: 0.2669317126274109 | loss_real: 0.0831474323272705 | loss_fake: 0.45074865221977234
[Epoch 12/20] | D loss: 3.89742043939277 | D loss: 0.0524422017620078 | loss_real: 0.0807244078731154 | loss_fake: 0.04891379612107801
[Epoch 13/20] | D loss: 4.70976226464697 | D loss: 0.1778805914974003 | loss_real: 0.1380549615433502 | loss_fake: 0.08548658470275961
[Epoch 14/20] | D loss: 7.0576677322387695 | D loss: 0.0890472521781821 | loss_real: 0.0340731731057167 | loss_fake: 0.14560213685835706
[Epoch 15/20] | D loss: 12.151673073131614 | D loss: 0.502876468022235 | loss_real: 0.0065896474445417 | loss_fake: 1.1762163361576538
[Epoch 16/20] | D loss: 10.298261614245605 | D loss: 0.323848263620256775 | loss_real: 0.0043887184001505375 | loss_fake: 0.6429768280503351
[Epoch 17/20] | D loss: 4.76896167987628 | D loss: 0.1106170497212219 | loss_real: 0.0512398736658896 | loss_fake: 0.1698938342824402
[Epoch 18/20] | D loss: 4.70938662505152 | D loss: 0.1778805914974003 | loss_real: 0.013681958989919707 | loss_fake: 0.3414602682324988
[Epoch 19/20] | D loss: 7.244987487792369 | D loss: 0.07223927229642688 | loss_real: 0.038334932178258896 | loss_fake: 0.10614361613988876
[Epoch 20/20] | D loss: 8.073629379272461 | D loss: 0.12738255202770233 | loss_real: 0.02396019990324974 | loss_fake: 0.24178451521802621
```

```
In [16]: fig = plt.figure(figsize=(15,10))
x_epochs = list(range(1, epochs+1))
plt.plot(x_epochs, G_losses, 'b-o', x_epochs, D_losses, 'r-x')
plt.legend(['Generator Loss', 'Discriminator Loss'])
plt.title('Optimizer Loss vs Epoch')
plt.xlabel('Epoch')
plt.ylabel('Loss')
```



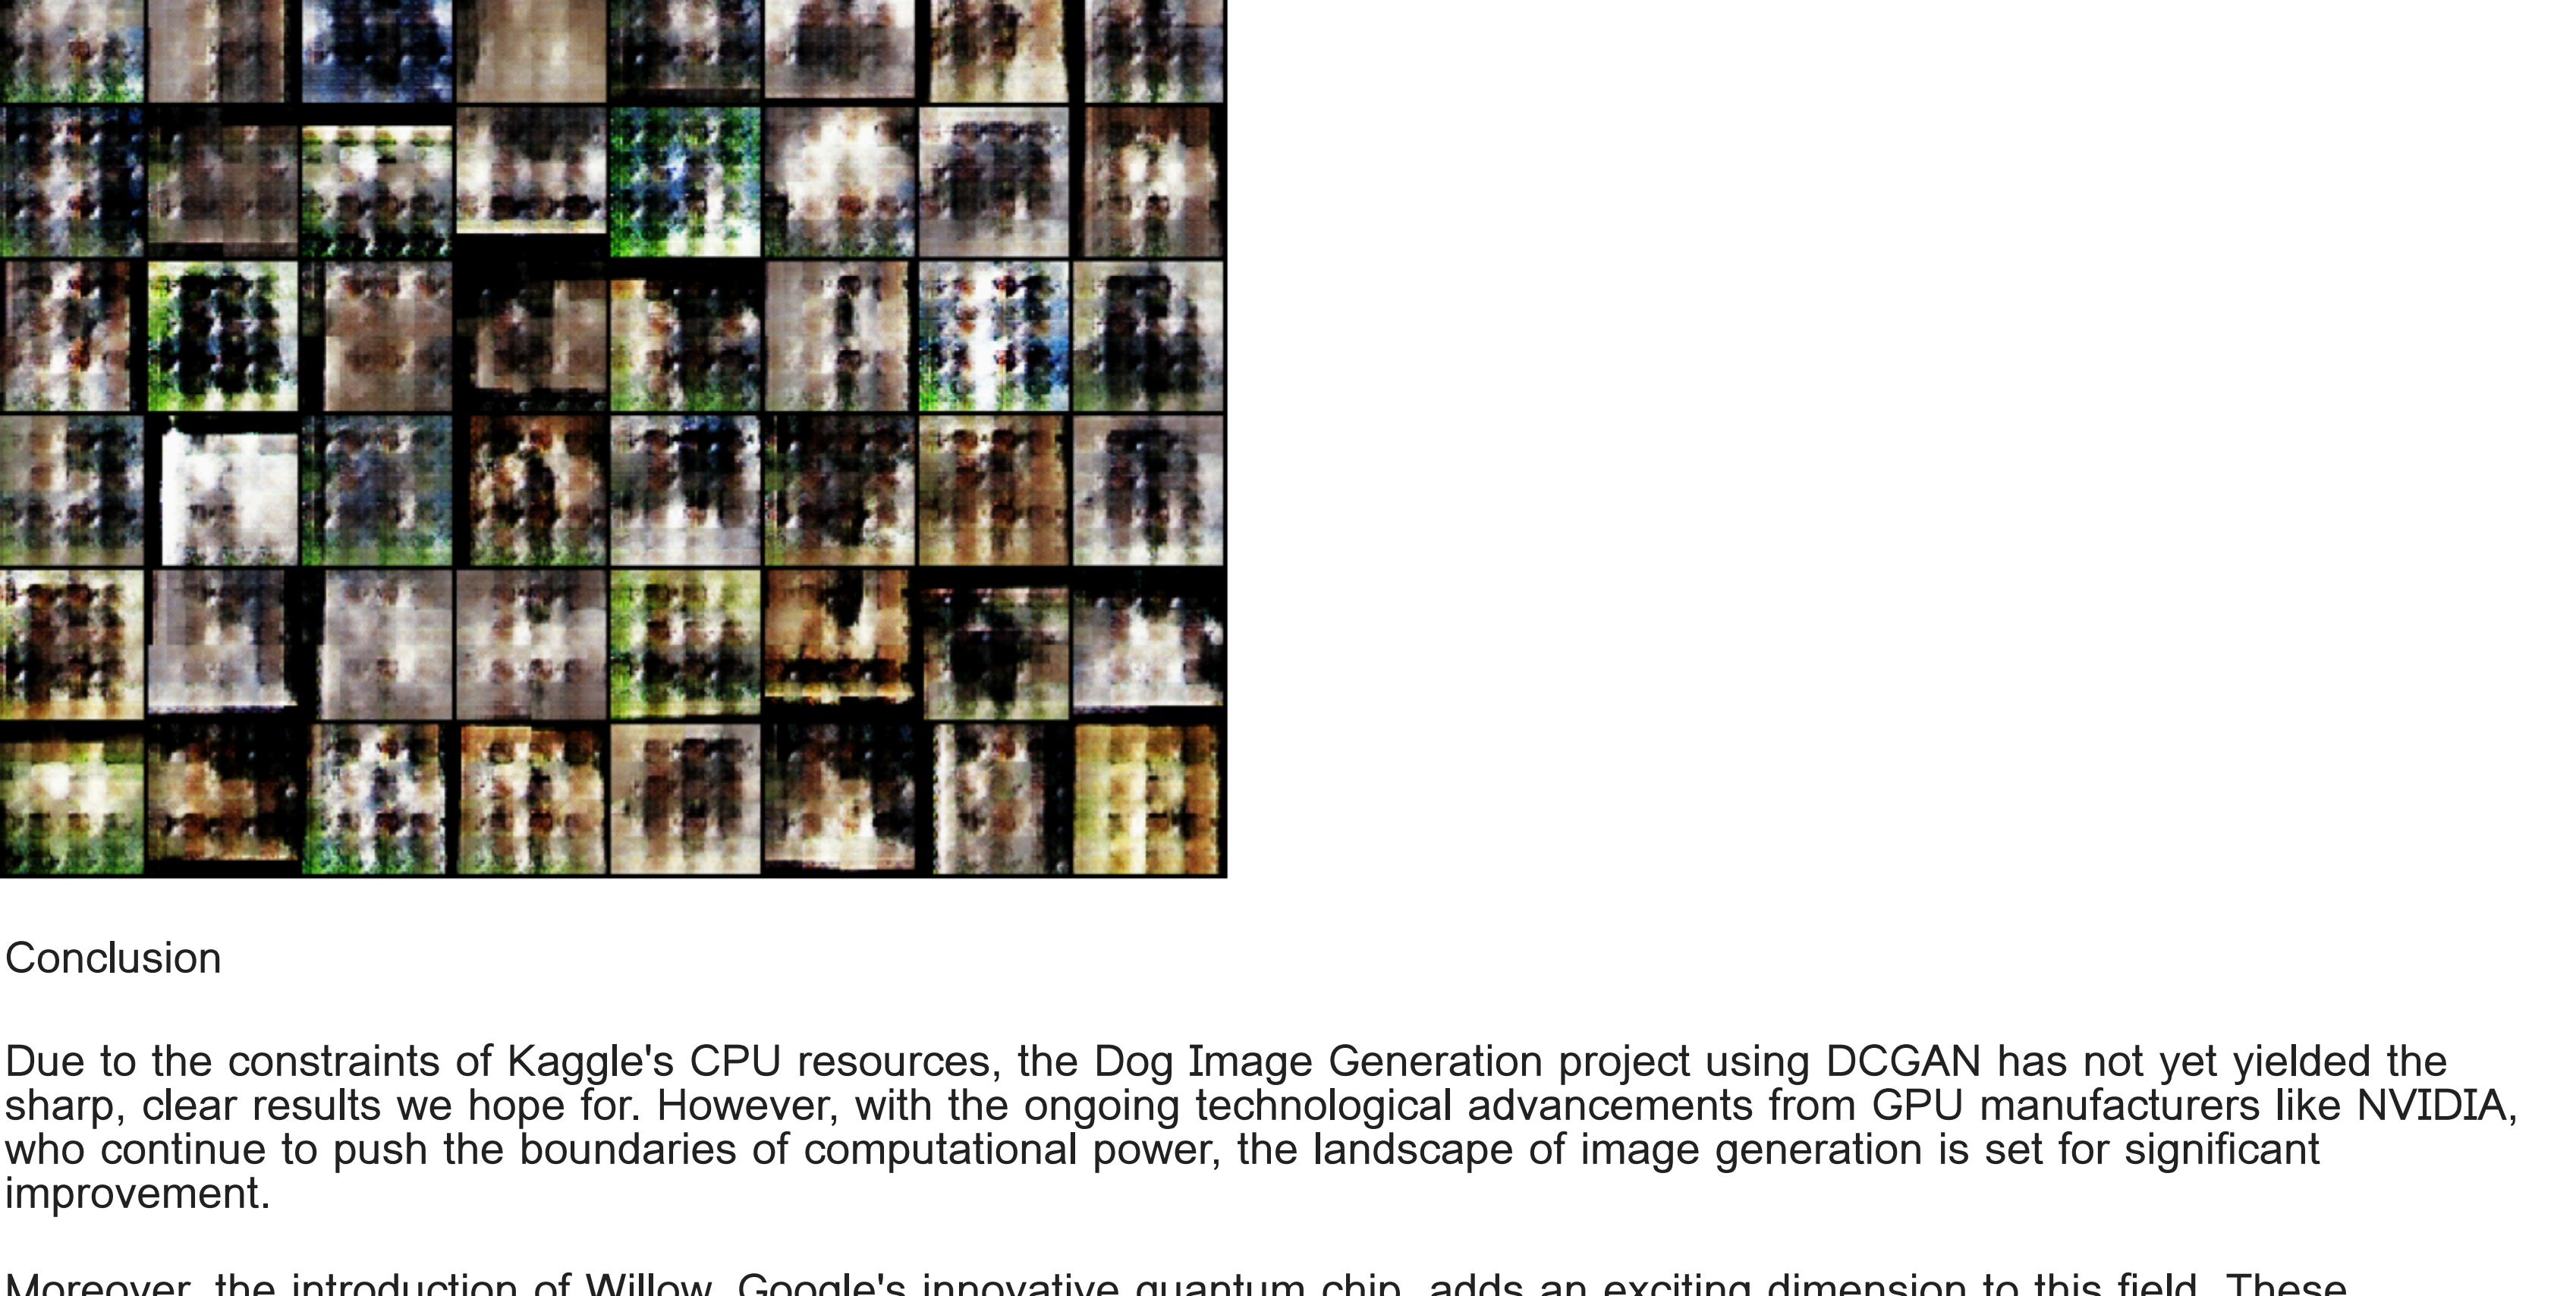
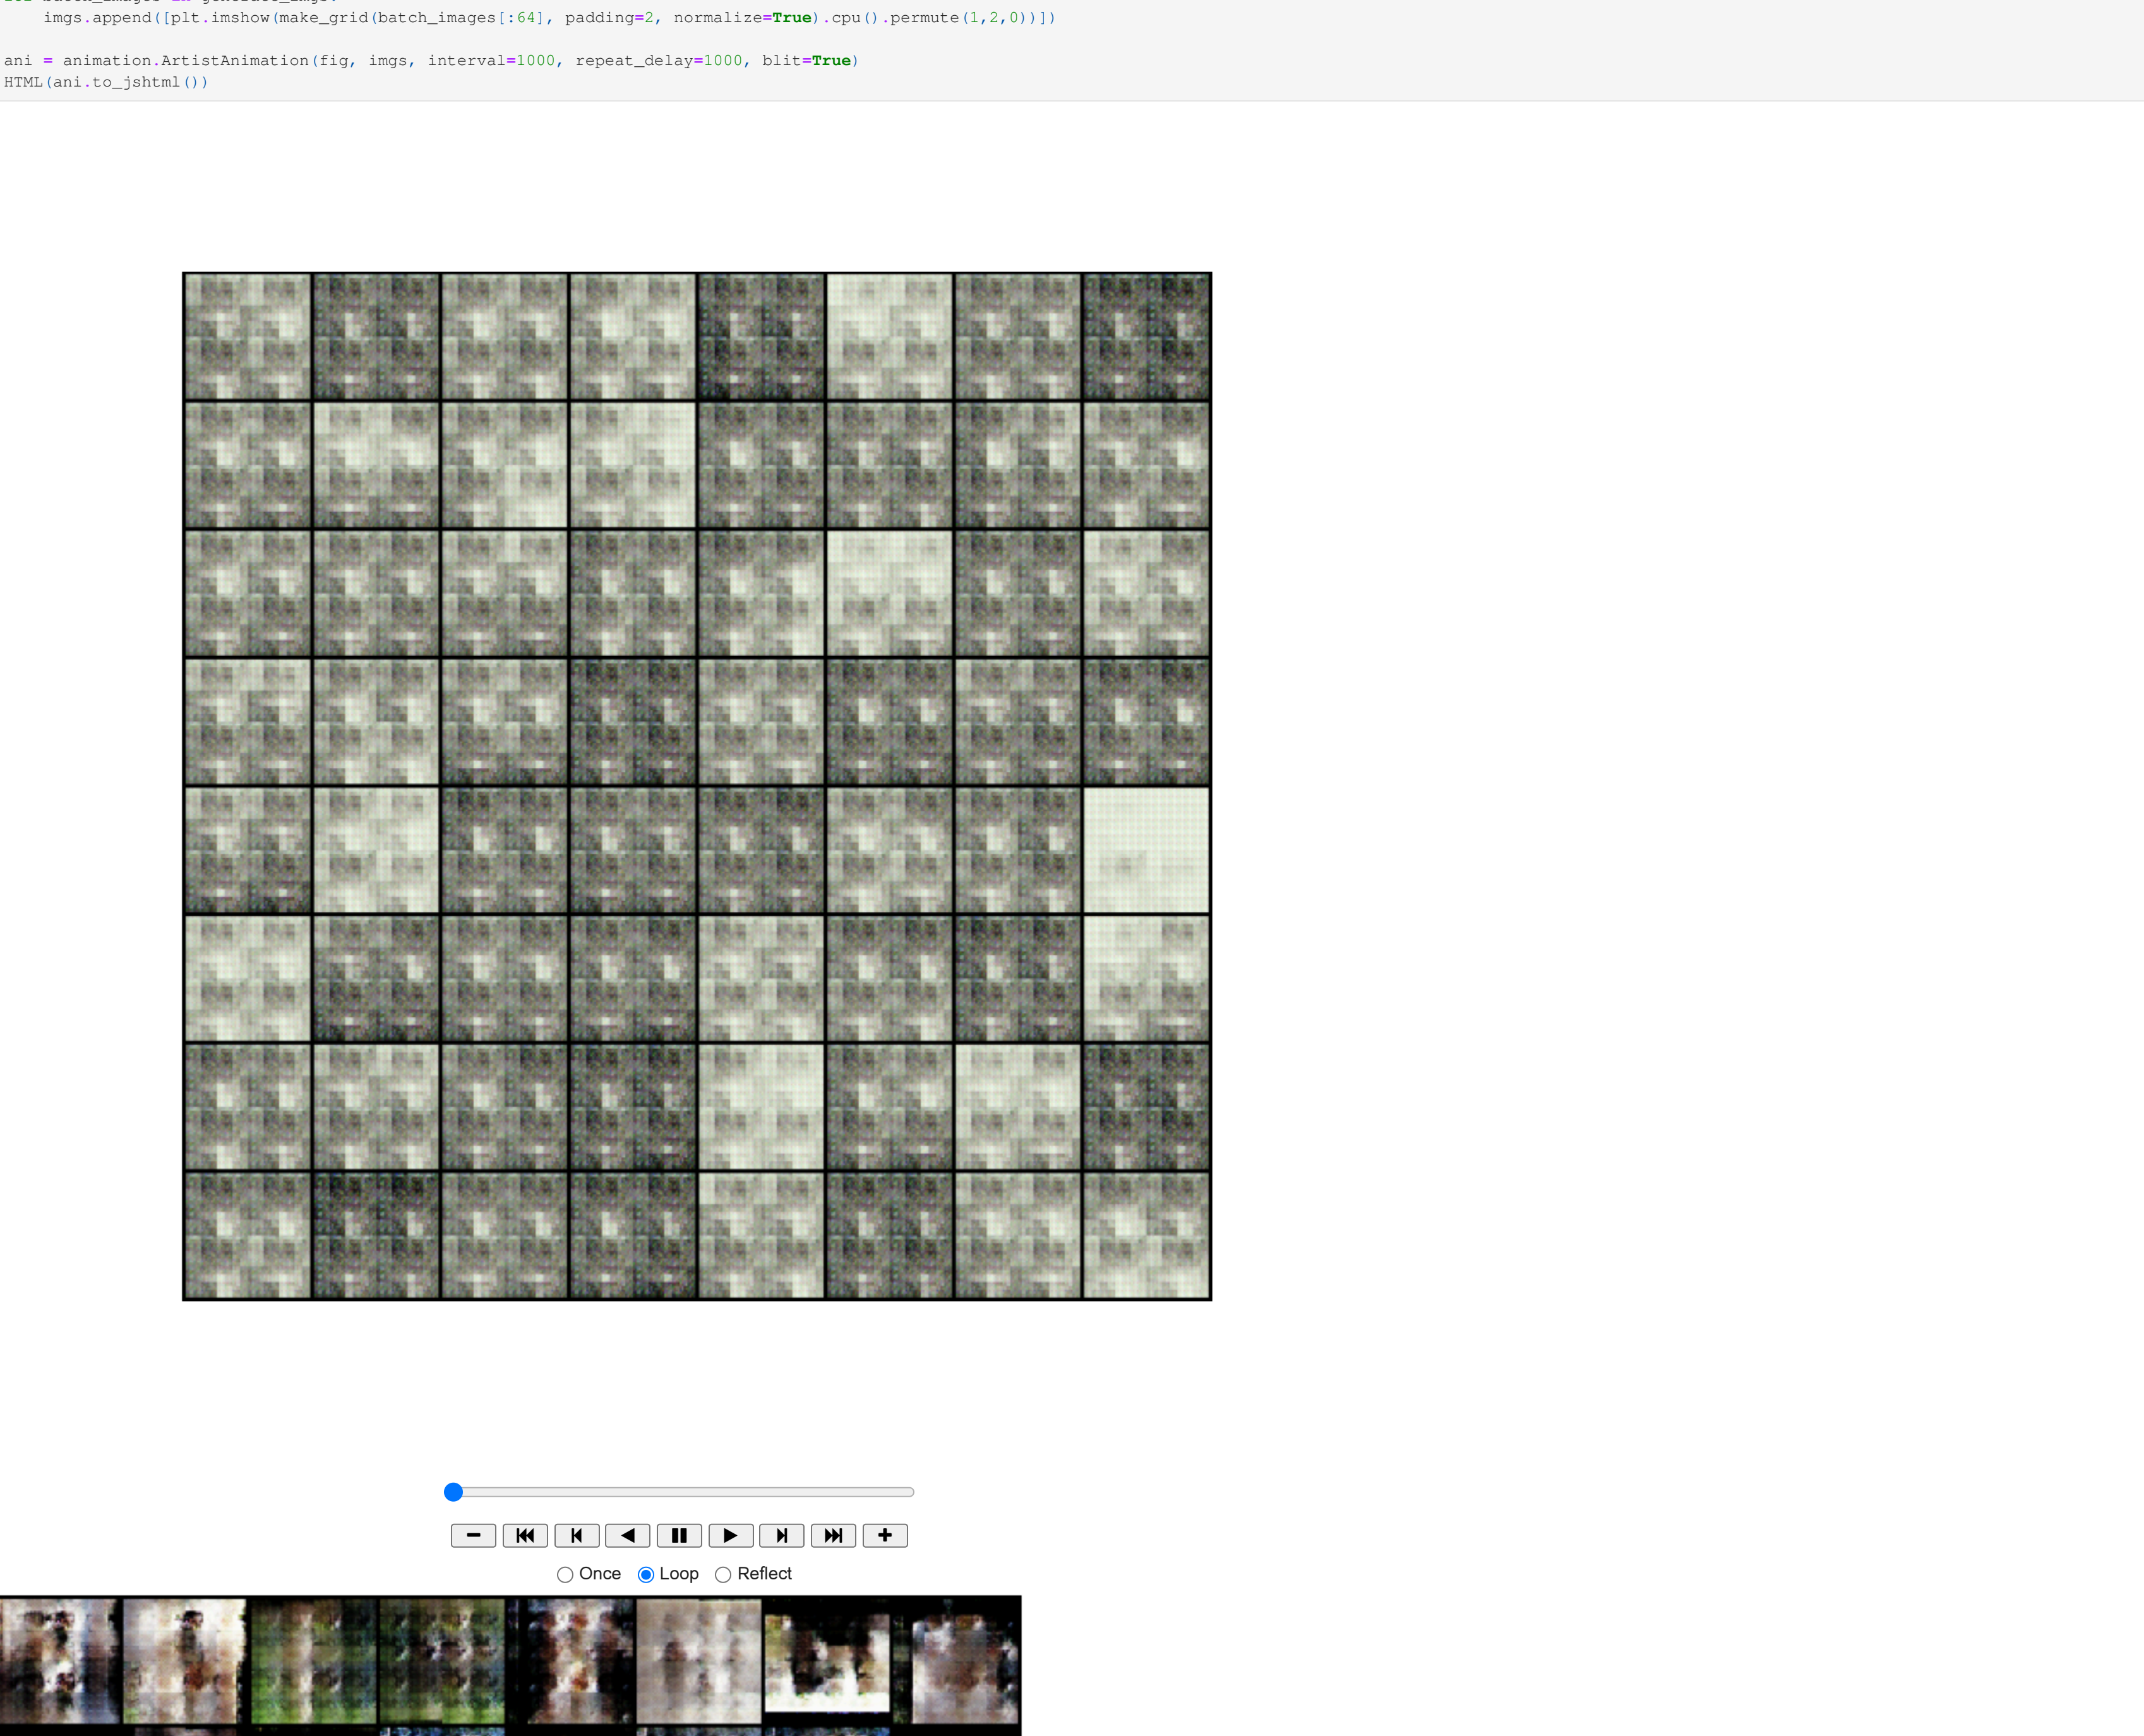
```
In [17]: import matplotlib.animation as animation
from IPython.display import HTML

fig = plt.figure(figsize=(15,15))
plt.axis('off')

imgs = []
for batch_images in generate_imgs:
    imgs.append(plt.imshow(make_grid(batch_images[64], padding=2, normalize=True).cpu().permute(1,2,0)))

ani = animation.ArtistAnimation(fig, imgs, interval=1000, repeat_delay=1000, blit=True)
HTML(ani.to_html())

Out[17]:
```



## Conclusion

Due to the constraints of Kaggle's CPU resources, the Dog Image Generation project using DCGAN has not yet yielded the sharp, clear results we hope for. However, with the ongoing technological advancements from GPU manufacturers like NVIDIA, who continue to push the boundaries of computational power, the landscape of image generation is set for significant improvement.

Moreover, the introduction of Willow, Google's innovative quantum chip, adds an exciting dimension to this field. These developments suggest a bright future for generative AI in image creation. Indeed, companies like xAI are already capitalizing on these technological leaps by offering rapid, high-quality image generation services to the public at no cost, responding to user prompts with remarkable efficiency and detail.